

TECHNICAL DOCUMENTATION SUITE

Software Architecture Document

Decoupled Architecture, Threat Modeling & Multi-Standard Engine

Platform Version:	2.5.0 Enterprise Edition
Compliance Scope:	OWASP AISVS v1.0, ISO/IEC 42001:2023, Google Cloud Security
License:	Apache License 2.0 (Open Source)

AEGIS AI ASSURANCE PLATFORM

Software Architecture Document (SAD)

Version: 2.5.0 Enterprise (2026 Edition)

Documentation Standard: Based on IEEE 1471 / ISO/IEC 42010 and the 4+1 View Model

Classification: Technical Architecture and Design Document for Principal Engineers and Auditors

1. Introduction and Architectural Goals

1.1 Purpose

This document formally describes the software architecture of **AegisAI**, a mission-critical platform focused on quality assurance, cybersecurity, and regulatory compliance in Generative Artificial Intelligence and Machine Learning systems.

1.2 Core Quality Attributes

- Deterministic Hyper-Factuality (QA-01):** Zero tolerance for hallucinations. The AI must evaluate exclusively based on the provided documentary evidence, rejecting unsupported inferences.
- Complete Output Capacity Without Truncation (QA-02):** The ability to audit over 40 comprehensive technical controls without exceeding the strict 8,192-token output limit per LLM call.
- Immutability and Forensic Traceability (QA-03):** Compliance with OWASP AI-SVS Invariant 4: audit logs cannot be altered or deleted once adjudicated.
- Resiliency and High Availability (QA-04):** Fault tolerance against transient network drops, cloud provider quota limits, or model saturation.
- Performance and Low Latency (QA-05):** Legal framework ingestion and real-time processing with sub-second latencies on structured tasks.

2. Design Patterns and Architectural Decisions

2.1 Map-Reduce Pattern for Distributed Inference

- Problem:** Contemporary LLM models feature massive input windows (1M to 2M tokens), but maintain a rigid generation limit (**completion tokens**) of 8,192 tokens. Attempting to audit 12 chapters (44 controls) with complete mitigation recipes in a single request causes the model to severely compress or truncate its responses.
- Solution:** A Map-Reduce architecture was implemented in `server.ts` :
- MAP Phase:** Standard chapters are partitioned into 3 or 4 parallel batches (e.g., Batch 1: C01-C03; Batch 2: C04-C06; Batch 3: C07-C09; Batch 4: C10-C12). Each batch is sent independently to the LLM with the strict structured schema.
- REDUCE Phase:** The orchestrator unifies the result matrices in memory, calculates the overall weighted score, and consolidates non-conformities into the final `AuditReport` object.

2.2 Server-Sent Events (SSE) Streaming Pattern

- **Problem:** A comprehensive technical evaluation of over 40 controls takes between 20 and 45 seconds. A traditional synchronous HTTP call would cause UI freezes or proxy timeouts.
- **Solution:** The `/api/evaluate` route operates under the SSE protocol (`text/event-stream`), dispatching granular progress events (`{ type: 'progress', progress: 50, message: '...' }`) and keeping the connection alive with `: keep-alive\n\n` heartbeats every 12 seconds.

2.3 Resilient Dual Persistence Pattern

- **Problem:** Relying exclusively on a cloud database (Cloud Firestore) exposes the application to outages caused by security rule propagations, disconnections, or permission locks.
- **Solution:** Every regulatory framework or audit is immediately stored in the local file system (`config/frameworks/{id}.json`) and replicated asynchronously and non-blockingly to Firestore. If Firestore returns a `PERMISSION_DENIED` or a network failure occurs, the system continues to operate at 100% capacity from the local cache.

2.4 Shared Responsibility Deduplication Pattern

- **Problem:** In architectures deployed on Google Gemini Enterprise SaaS, penalizing the client for failing to prove how model weights are encrypted on physical TPU hardware is an audit error that misdirects attention away from application security.
- **Solution:** The evaluation engine incorporates automated logic to filter out infrastructure provider responsibilities (such as physical hardware security, cloud-level data centers, and multi-tenant isolation), focusing exclusively on the customer-controlled application layer (prompts, RAG, filters, API security, and governance).

2.5 Multi-Provider Decoupling and Universal Inference Pattern

- **Problem:** Forcing an exclusive dependency on a single AI provider (vendor lock-in) prevents corporate clients from utilizing their existing agreements with OpenRouter, OpenAI, or local inference servers (vLLM, Ollama, LiteLLM).
- **Solution:** A decoupled engine centered around `executeAIWithFallback` was implemented:
- If the configured provider is `gemini`, it utilizes the `@google/genai` SDK with its resiliency hierarchy (`gemini-3.8-flash -> gemini-3.1-flash-lite -> gemini-3.7-flash`).
- If the provider is `openrouter`, `openai`, or `custom_openai_compatible`, it utilizes a native HTTP client implementing the universal `/v1/chat/completions` specification with 120s timeouts, exponential backoff, strict JSON extraction, and input/output token accounting.

2.6 Remediated Architecture Generation Pattern (Drop-in Replacement)

- **Problem:** Providing a client solely with a report containing failing scores or non-conformity lists does not resolve the technical gap and creates operational friction.
- **Solution:** When failures or missing evidence are detected, the system orchestrates a synthetic remediation phase. The engine takes the user's original architecture and weaves in all missing safeguards (anti-injection sanitization, DLP anonymization, RAG isolation, immutable logs, and human scaling), producing a production-ready Markdown document (`.md` / `.txt`) that guarantees **100.0% regulatory compliance**.

2.7 Simultaneous Multi-Framework Orchestration Pattern (Compliance 360°)

- **Problem:** In production environments and regulated sectors (banking, retail, healthcare), an AI system must simultaneously satisfy multiple regulatory layers: technical cybersecurity (OWASP AI-SVS 1.0), AI governance and risk management (ISO/IEC 42001:2023), data privacy (LFPDPPP / INAI), and commercial transparency (LFPC / PROFECO). Executing fragmented audits multiplies latency and creates reporting silos.
- **Solution:** `runArchitectureEvaluation` supports the concurrent specification of multiple regulatory frameworks via `selectedStandards: string[ ]`. During the MAP phase, the engine aggregates batches from all selected standards,

injecting lineage metadata (`standardId` , `standardName` , `contextText`) into each inference block. In the REDUCE phase, the system calculates the global metric and structures the `standardsBreakdown` array, where each standard preserves its independent score (0-100%), total evaluated controls, and compliance status, allowing for the rendering of Compliance 360° dashboards and the export of consolidated executive reports.

2.8 Live Cloud Infrastructure Inspection Pattern (GCP Live Scanner - Zero-Footprint)

- **Problem:** Auditing exclusively static documents or written diagrams introduces human bias ("wishful thinking" or discrepancies between specification and actual implementation). Likewise, deploying software agents or functions within the client's project generates friction with corporate cybersecurity areas due to cost, stability risks, and approval processes.
- **Solution:** The `gcpScannerService` was designed to interact exclusively with the **Google Cloud Control Plane (Management APIs)** via read-only `HTTP GET` queries. The service supports in-memory authentication via ADC (Cloud Run), temporary Bearer tokens, or Service Accounts without persisting keys to disk. It inspects critical configurations of Cloud IAM, Vertex AI, Cloud Run, Cloud Storage, and Secret Manager, synthesizing a structured technical manifest () that feeds directly into the multi-normative Map-Reduce engine, guaranteeing a live cloud-reality-based audit with **zero footprint and zero resources created on the client side**.

3. Logical Component View

flowchart TD

```

subgraph ClientLayer ["Client Layer (SPA Frontend)"]
    UI["React 19 + Tailwind CSS v4\nView-based Code-Splitting"]
    SSE_Client["SSE Event Consumer\nReactive Streaming"]
    PDF_Engine["jsPDF + AutoTable\nDynamic On-Demand Loading"]
    RemediationUI["Remediated Architecture\nPreview & Download Modal"]
end

subgraph GatewayLayer ["Routing & Security Layer (Express)"]
    SecHeaders["OWASP Headers\n(nosniff, SAMEORIGIN, PermPolicy)"]
    DotfileBlocker["Anti-Probing Filter\n(/.env, .key, .git Blocking)"]
    AuthMiddleware["RBAC Authentication\n(x-admin-token, x-api-key)"]
end

subgraph CoreEngine ["Multi-Provider AI Evaluation Core"]
    Orchestrator["runArchitectureEvaluation\n(Map-Reduce Orchestrator)"]
    Decomposer["decomposeLegalDocument\n(Universal BYOF Engine)"]
    Remediator["remediatedArchitectureGenerator\n(100% Compliance Synthesis)"]
    Dispatcher["executeAIWithFallback\n(Central Dispatcher)"]
    GeminiClient["Google GenAI SDK\n(503/429 Fallback Chain)"]
    UniversalOpenAI["Universal /v1/chat/completions Client\n(OpenRouter / OpenAI / vLLM / Ollama)"]
end

subgraph PersistenceLayer ["Dual Persistence Layer"]
    LocalFS["Local File System\nconfig/frameworks/*.json & admin-settings.json"]
    CloudFirestore["Google Cloud Firestore\nCollections: audits, custom_frameworks"]
end

UI --> SecHeaders
SecHeaders --> DotfileBlocker
DotfileBlocker --> AuthMiddleware
AuthMiddleware --> Orchestrator
AuthMiddleware --> Decomposer
AuthMiddleware --> Remediator
Orchestrator --> Dispatcher
Decomposer --> Dispatcher
Remediator --> Dispatcher
Dispatcher -->|gemini| GeminiClient
Dispatcher -->|openrouter/openai/local| UniversalOpenAI
Orchestrator --> LocalFS
Orchestrator --> CloudFirestore
Decomposer --> LocalFS
Decomposer --> CloudFirestore
Orchestrator -. SSE Events .-> SSE_Client
Remediator -. .md Specification .-> RemediationUI

```

4. Domain and Data Models

4.1 `AuditReport` Entity

Representative structure of the issued technical report:

```

interface AuditReport {
  id: string; // Unique identifier (e.g., LFPDPPP-2026-0001)
  date: string; // ISO 8601 timestamp
  systemName: string; // Name of the audited system
  technicalLead: string; // Responsible technical lead
  email: string; // Contact email
  description: string; // Evaluated technical specification
  standard: 'AI-SVS' | 'ISO-42001' | 'FULL' | 'CUSTOM';
  customFrameworkId?: string; // Custom law ID if standard === 'CUSTOM'
  customFrameworkName?: string; // Descriptive name of the regulation
  targetLevel: 'L1' | 'L2' | 'L3' | 'AUTO';
  assessedLevel: 'L1' | 'L2' | 'L3';
  levelAssessmentRationale: string; // Technical justification for the level
  techPlatform: string; // gemini-enterprise, vertex-ai, self-hosted
  techPlatformName: string; // Commercial name of the platform
  deploymentType: 'SaaS' | 'PaaS' | 'IaaS' | 'Self-Hosted';
  overallScore: number; // Weighted score (0-100%)
  executiveSummary: string; // Executive summary of risks
  chapters: StandardChapter[]; // Evaluated chapters with their tested points
  modelUsed: string; // Exact Gemini model used
  latencyMs: number; // Total inference time
  promptTokens: number; // Input tokens
  completionTokens: number; // Generated output tokens
}

```

4.2 `PuntoProbado` Entity and 4D Mitigation Recipe

```

interface PuntoProbado {
  id: string; // Control identifier (e.g., LFPDPPP-03.1)
  name: string; // Requirement title
  status: 'Aprobado' | 'Requiere Revisión' | 'Falta Evidencia' | 'No Aplica';
  findings: string; // Findings and technical justification
  fuenteEvidencia?: string; // Exact textual quote from the documentation
  isInheritedFromPlatform?: boolean; // True if covered by Google Cloud SaaS
  inheritedPlatformName?: string; // Name of the accredited cloud
  level?: 'L1' | 'L2' | 'L3'; // Verification level
  isMandatoryForTargetLevel?: boolean;
  remediationDetails?: {
    strategy: string; // Dimension 1: Architectural strategy
    actionableSteps: string[]; // Dimension 2: Step-by-step checklist
    codeOrConfigExample: string; // Dimension 3: Real code or manifest
    verificationRecipe: string; // Dimension 4: Pentesting script / curl QA
    effortLevel: string; // Estimated effort (< 1 day, 1-3 days, etc.)
    recommendedTools: string[]; // Recommended libraries (Presidio, Sigstore)
  };
}

```

5. Process View: Sequence Diagrams

5.1 Audit Flow with Map-Reduce and SSE Streaming

```
sequenceDiagram
    autonumber
    actor User as Auditor / User
    participant FE as React Frontend
    participant API as Express Backend (/api/evaluate)
    participant Engine as runArchitectureEvaluation
    participant Gemini as Google Gemini AI Studio
    participant Storage as Dual Persistence

    User->>FE: Submits Audit Form (Specification + Platform)
    FE->>API: POST /api/evaluate (SSE Handshake)
    API-->>FE: HTTP 200 (text/event-stream connected)
    API->>Engine: Initiates Evaluation

    loop For each Map Batch (Batches 1 to 4)
        Engine->>Gemini: generateContent(Batch Prompt + Schema + Shared Responsibility)
        Gemini-->>Engine: Structured JSON with Controls and Textual Quotes
        Engine-->>FE: SSE Event (progress: XX%, batch completed)
    end

    Engine->>Engine: REDUCE (Consolidation of scores and findings)
    Engine->>Storage: Saves to local config/frameworks/ and Cloud Firestore
    Engine-->>FE: SSE Event (type: 'result', data: AuditReport)
    FE->>User: Renders Report with Inheritance Badges and 4D Recipes
```

6. Threat Modeling and Security

AegisAI has been subjected to and hardened against the 10 threat categories of the **OWASP Top 10 for LLM Applications**:

Identifier	Threat Vector	Defense Mechanism Implemented in AegisAI
LLM01	Prompt Injection	Strict delimiters <code><user_candidate_architecture></code> , deterministic parameters (<code>temp: 0.0</code>), and isolation of input data as untrusted.
LLM02	Insecure Output Handling	Rigid typing in JSON schemas, anti-XSS sanitization in React, and secure rendering in PDF and Markdown exporters.
LLM03	Data Poisoning	Mandatory structural validation during regulatory ingestion (<code>frameworkDecomposerSchema</code>).
LLM04	Model Denial of Service	Maximum payload limit in Express (15MB), HTTP 413 rejection for larger payloads, and SSE streaming with flow control.
LLM05	Supply Chain Vulnerabilities	Minimized dependencies, lockfile enforcement via <code>package-lock.json</code> , and zero dependencies with critical CVEs.
LLM06	Information Disclosure	Automatic log masking that prevents the exfiltration of API keys (<code>AlzaSy...</code>) or tokens in error responses.
LLM07	Insecure Plugin Design	Strict middlewares <code>authenticateAdmin</code> (<code>x-admin-token</code>) and <code>authenticateApi</code> (<code>x-api-key</code>) for sensitive endpoints.
LLM08	Excessive Agency	OWASP AI-SVS Invariant 4: Audits are immutable; no HTTP DELETE routes exist to remove reports.
LLM09	Overreliance (Hallucination)	Mandatory textual evidence quotes; if the documentation does not mention it, the verdict is "Evidence Missing".
LLM10	Model Theft / File Probing	Anti-probing middleware that blocks requests to dotfiles (<code>/.env</code> , <code>.key</code> , <code>.git</code>), returning an immediate HTTP 404.

7. Cloud Deployment View

The AegisAI infrastructure architecture leverages the managed ecosystem of **Google Cloud Platform (GCP)** to ensure high availability (99.95%), zero server maintenance, and strict network isolation:

```
flowchart LR
    subgraph ClientPerimeter ["Client / Auditor Perimeter"]
        Browser["Secure Browser\n(Chrome / Safari / Firefox)"]
    end

    subgraph GCP_Project ["Google Cloud Project (Enterprise / Public Tenant)"]
        subgraph IngressLayer ["Ingress and CDN Layer"]
            CloudRunFront["Google Cloud Run Ingress\n(HTTPS TLS 1.3 / Port 8080)"]
            FirebaseAuth["Firebase Auth / Identity Platform\n(Configurable OAuth Filter: VITE_ALLOW..."]
        end

        subgraph ComputeLayer ["Serverless Compute Layer"]
            Container["AegisAI Container\n(Node.js 22 LTS, Non-root Runner: aegisuser)\nExpress API"]
            ServiceAccount["Runtime SA: aegis-runner-sa\nRoles: aiplatform.user, datastore.user, vie..."]
        end

        subgraph DataAndAILayer ["Managed Persistence and AI Services"]
            FirestoreDB["Cloud Firestore (Native Mode)\nCryptographic Rules firestore.rules"]
            VertexAI["Google Cloud Vertex AI\nGemini 2.5 / 3.x Flash & Pro Models"]
            ArtifactReg["Artifact Registry (Docker)\ncloud-run-source-deploy"]
        end
    end

    Browser -->|HTTPS GET/POST| CloudRunFront
    Browser -. ->|OAuth Popup| FirebaseAuth
    CloudRunFront --> Container
    Container --- ServiceAccount
    ServiceAccount -->|IAM Read/Write| FirestoreDB
    ServiceAccount -->|Inference and Telemetry| VertexAI
    Container -. ->|Image Deployment| ArtifactReg
```

7.1 Infrastructure Components and Service Mapping

GCP Component	Technical Service	Role and Responsibility
Serverless Compute	Google Cloud Run	Hosts the optimized monolithic container (`dist/server.cjs` + static SPA). Scales from 0 to 10 instances based on demand.
Container Registry	Google Artifact Registry	Stores and scans signed multi-stage Docker images.
Identity and Access	Firebase Auth / Identity Platform	Provides federated corporate sign-in with Google Workspace and validates membership in authorized corporate domains.
NoSQL Database	Cloud Firestore (Native)	Stores issued reports, audit metadata, and custom regulatory catalogs with distributed persistence.
Inference and AI	Vertex AI / Gemini API	Executes audit Map-Reduce queries and extracts factual project telemetry.
CI/CD Automation	Google Cloud Build	Orchestrates cloud-based builds without requiring Docker on developer workstations.